

Kubernetes (GKE) Practical Lab Guide

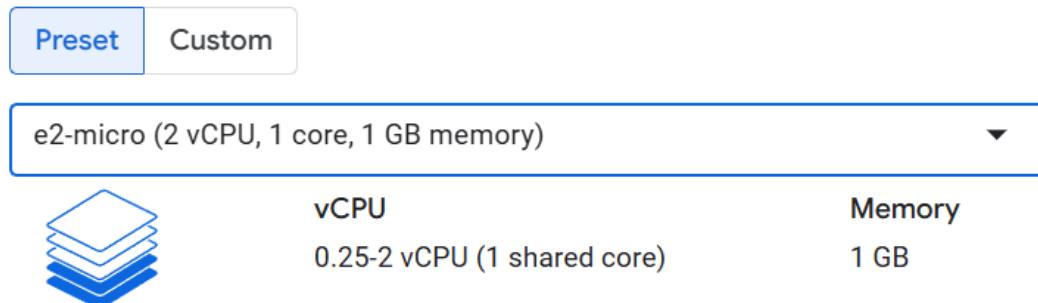
✓ Task 1 – Create a GKE Cluster & Connect via Console

Steps 1 To Create GKE Cluster

1. Go to **Kubernetes Engine > Clusters**.
2. Click “**Create**” > Choose “**Standard**”.
3. Configure the following:
 - **Cluster Name:** my-Cluster1
 - **Location Type:** Zonal (e.g., us-central1-a)
 - **Number of Nodes:** 3

Machine type

Choose a machine type with preset amounts of vCPUs and memory that suit most workloads. Or, you can create a custom machine for your workload's particular needs. [Learn more](#) ↗

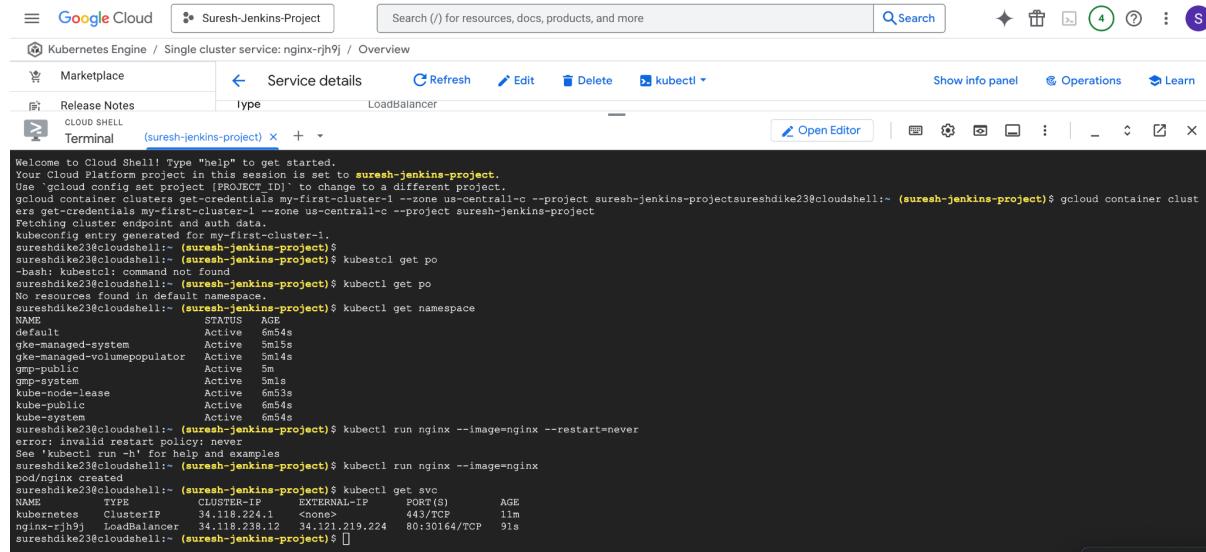


4. Click “**Create**” and wait for the cluster to be provisioned.

The image shows the Google Cloud Kubernetes Engine Clusters page. The left sidebar shows 'Clusters' selected. The main area displays the 'my-first-cluster-1' cluster details. It includes sections for 'Overview', 'Health' (100% healthy), 'Upgrade' (100% up to date), 'Estimated monthly cost' (₹0.00 / month), and a table of cluster metrics. The table has columns: Status, Name, Location, Tier, Number of nodes, Total vCPUs, Total memory, Notifications, and Labels. The 'my-first-cluster-1' row is selected, showing 'Status' as 'Green' (healthy), 'Name' as 'my-first-cluster-1', 'Location' as 'us-central1-c', 'Tier' as 'Standard', 'Number of nodes' as 3, 'Total vCPUs' as 3, 'Total memory' as 5.1 GB, 'Notifications' as none, and 'Labels' as none.

Step 2: Connect to GKE Cluster

- Once the cluster is ready, click on its name.
- Click “**Connect**” > “**Run in Cloud Shell**”.
- This will automatically run the following command:
gcloud container clusters get-credentials my-Cluster1 --zone us-central1-a



```
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to suresh-jenkins-project.
Use "gcloud config set project [PROJECT_ID]" to change to a different project.
gcloud container clusters get-credentials my-Cluster1 --zone us-central1-a --project suresh-jenkins-project
Fetching Cluster endpoint and auth data
kubeconfig entry generated for my-Cluster1.
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$ kubectl get po
-bash: kubectl: command not found
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$ kubectl get po
No resources found
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$ kubectl get namespace
NAME          STATUS  AGE
default       Active  6m54s
gke-managed-system  Active  5m15s
gke-managed-volumepopulator  Active  5m14s
gmp-public    Active  5m
gmp-system    Active  5m15s
kube-dns      Active  6m33s
kube-public   Active  6m54s
kube-system   Active  6m14s
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$ kubectl run nginx --image=nginx --restart=never
error: invalid restart policy: never
See 'kubectl run -h' for help and examples
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$ kubectl run nginx --image=nginx
pod/nginx created
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$ kubectl get svc
NAME          TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
kubernetes   ClusterIP  34.118.224.1  <none>        443/TCP      11m
nginx-rjh9j  LoadBalancer  34.118.238.12  34.121.219.224  80:30164/TCP  91s
sureshdk2e36@cloudshell:~ (suresh-jenkins-project)$
```

✓ Task 2 – Deploy Nginx Pod & Expose It on Port 80

Step 1: Create Nginx Pod

kubectl run nginx --image=nginx --restart=Never

Step 2: Expose the Pod using LoadBalancer

kubectl expose pod nginx-pod --type=LoadBalancer --port=80

Step 3: Get External IP

kubectl get svc

- Wait for EXTERNAL-IP to be assigned (may take 1-2 mins).
- Access your Nginx app using:
<http://<EXTERNAL-IP>>

CLOUD SHELL Terminal (suresh-jenkins-project) + Open Edit

```

-bash: kubectl: command not found
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl get po
No resources found in default namespace.
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl get namespace
NAME          STATUS  AGE
default       Active  6m54s
gke-managed-system  Active  5m15s
gke-managed-volumepopulator  Active  5m14s
gmp-public    Active  5m
gmp-system    Active  5m1s
kube-node-lease  Active  6m53s
kube-public   Active  6m54s
kube-system   Active  6m54s
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl run nginx --image=nginx --restart=never
error: invalid restart policy: never
See 'kubectl run -h' for help and examples
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl run nginx --image=nginx
pod/nginx created
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  34.118.224.1  <none>        443/TCP      11m
nginx-rjh9j  LoadBalancer  34.118.238.12  34.121.219.224  80:30164/TCP  91s
sureshdike23@cloudshell:~ (suresh-jenkins-project)$
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl run jenkins-pod --image=jenkins/jenkins:lts --restart=Never
pod/jenkins-pod created
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  34.118.224.1  <none>        443/TCP      14m
nginx-rjh9j  LoadBalancer  34.118.238.12  34.121.219.224  80:30164/TCP  3m48s
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
jenkins-pod-sn7xj  LoadBalancer  34.118.231.48  35.193.112.150  8080:31540/TCP  71s
kubernetes  ClusterIP  34.118.224.1  <none>        443/TCP      16m
nginx-rjh9j  LoadBalancer  34.118.238.12  34.121.219.224  80:30164/TCP  6m15s
sureshdike23@cloudshell:~ (suresh-jenkins-project)$

```

Google Cloud Suresh-Jenkins-Project Search (/) for resources, docs, products, and more

Kubernetes Engine / Workloads

Learn about Enterprise Workloads Refresh Deploy Create Job Delete

All Fleets

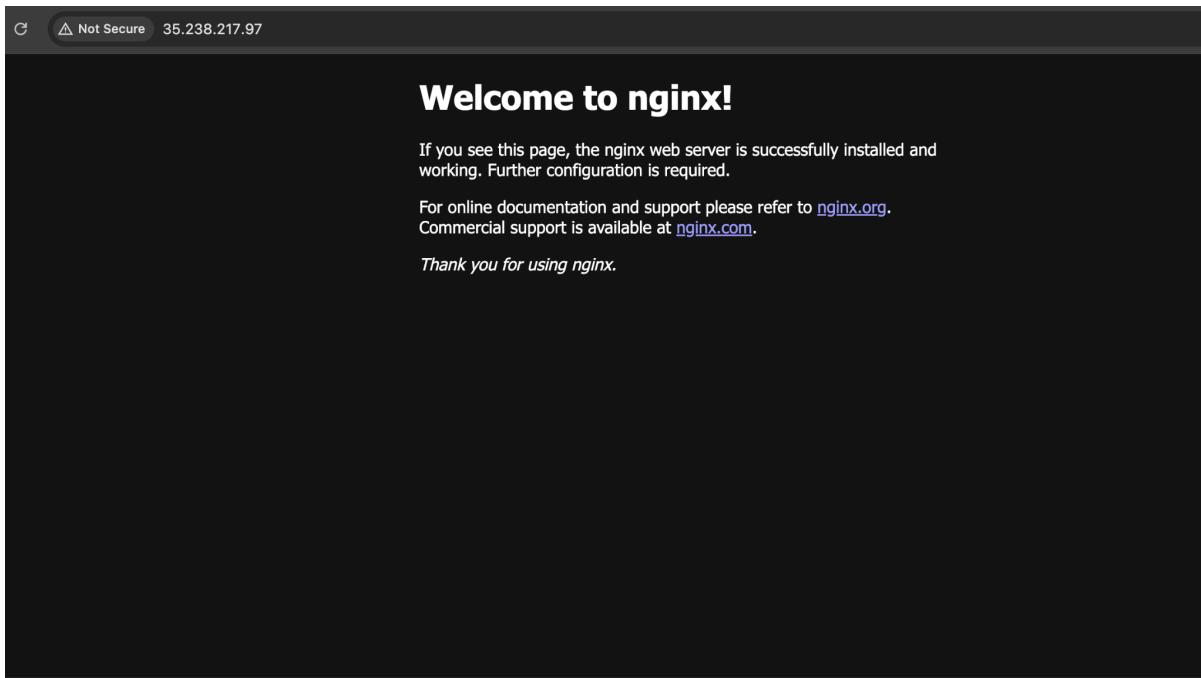
Resource Management

- Overview
- Clusters
- Workloads
- Teams
- Applications
- AI/ML New
- Secrets & ConfigMaps
- Storage
- Object Browser
- Upgrades New
- Backup for GKE
- Marketplace
- Release Notes

Overview Observability Cost Optimization

Filter Is system object : False Filter workloads

Name	Status	Type	Pods	Namespace	Cluster
jenkins-pod	Running	Pod	1/1	default	my-first-cluster-1
nginx	Running	Pod	1/1	default	my-first-cluster-1
nginx-deploy	OK	Deployment	3/3	nginx-namespace	my-first-cluster-1



✓ Task 3 – Deploy a Jenkins Pod & Expose It on Port 8080

Step 1: Deploy Jenkins Pod

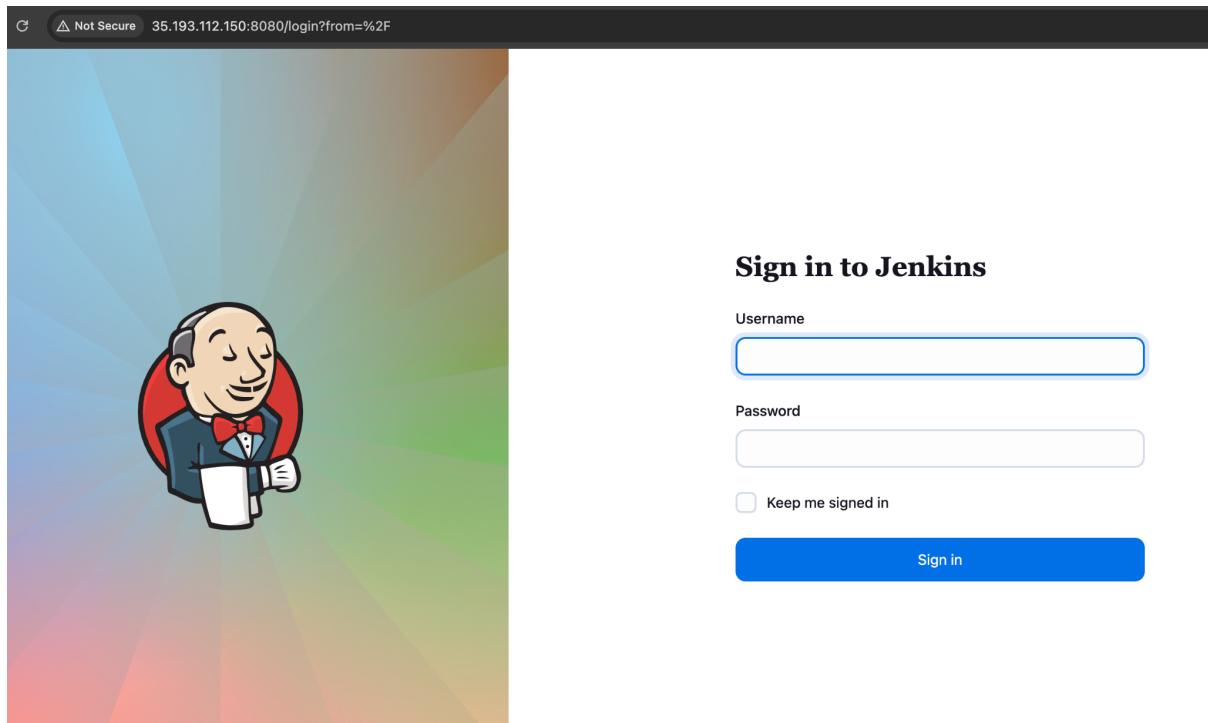
```
kubectl run jenkins --image=jenkins/jenkins:lts --restart=Never
```

Step 2: Expose the Pod as LoadBalancer

```
kubectl expose pod jenkins-pod --type=LoadBalancer --port=8080
```

Step 3: Access Jenkins UI

1. Run:
`kubectl get svc`
2. Get the **NodePort** and **External IP** of your cluster node.
3. Access Jenkins in the browser:
[`http://<EXTERNAL-IP>:<NODE-PORT>`](http://<EXTERNAL-IP>:<NODE-PORT>)



✓ Task 4 – Create an Nginx Deployment & Delete One Pod

Step 1: Create Nginx Deployment

```
kubectl create deployment nginx-deploy --image=nginx
```

Step 2: Check Running Pods

```
kubectl get pods
```

Step 3: Delete One Pod from the Deployment

Step 4: Observe Auto-Healing

```
kubectl get pods
```

- The deleted pod will be automatically recreated by the deployment controller.

```

sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl get pods -n nginx-namespace
NAME                         READY   STATUS    RESTARTS   AGE
nginx-deploy-c9d9f6c6c-7cd2n  1/1    Running   0          28m
nginx-deploy-c9d9f6c6c-dwt7m  1/1    Running   0          28m
nginx-deploy-c9d9f6c6c-jf6s5  1/1    Running   0          36m
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl delete pod nginx-deploy-c9d9f6c6c-7cd2n -n nginx-namespace
-bash: kubectl: command not found
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl delete pod nginx-deploy-c9d9f6c6c-7cd2n -n nginx-namespace
pod "nginx-deploy-c9d9f6c6c-7cd2n" deleted
sureshdike23@cloudshell:~ (suresh-jenkins-project)$ kubectl get pods -n nginx-namespace
NAME                         READY   STATUS    RESTARTS   AGE
nginx-deploy-c9d9f6c6c-5fsqg  1/1    Running   0          8s
nginx-deploy-c9d9f6c6c-dwt7m  1/1    Running   0          30m
nginx-deploy-c9d9f6c6c-jf6s5  1/1    Running   0          38m
sureshdike23@cloudshell:~ (suresh-jenkins-project)$

```

Google Cloud Suresh-Jenkins-Project

Kubernetes Engine / Deployment: nginx-deploy / Deployment overview

Deployment details

Cluster: my-first-cluster-1 Namespace: nginx-namespace Labels: app: nginx-deploy

Replicas: 3 updated, 3 ready, 3 available, 0 unavailable

Pod specification: Revision 1, containers: nginx

Horizontal Pod Autoscaler: Not configured

Vertical Pod Autoscaler: Not configured

Active revisions

Revision	Name	Status	Summary	Created on	Pods running/Pods total
1	nginx-deploy-c9d9f6c6c	OK	nginx: nginx	Apr 7, 2025, 11:07:59 PM	3/3

Managed pods

Revision	Name	Status	Restarts	Created on
1	nginx-deploy-c9d9f6c6c-jf6s5	Running	0	Apr 7, 2025, 11:08:00 PM
1	nginx-deploy-c9d9f6c6c-dwt7m	Running	0	Apr 7, 2025, 11:15:46 PM
1	nginx-deploy-c9d9f6c6c-5fsqg	Running	0	Apr 7, 2025, 11:46:22 PM

Exposing services

Name	Type	Endpoints
nginx-deploy-service	Load balancer	35.238.217.97:80

Google Cloud Suresh-Jenkins-Project

Kubernetes Engine Pod details

nginx-deploy-c9d9f6c6c-jf6s5

Details Events Logs App Errors (0) YAML

Select the Cloud Monitoring account to see charts.

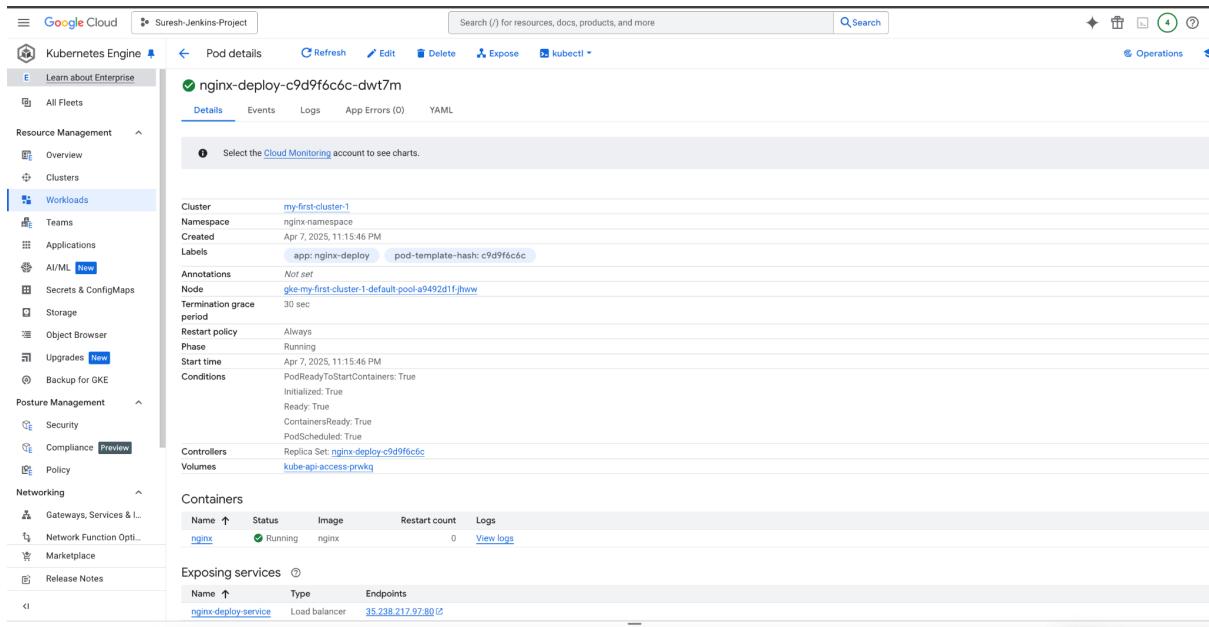
Cluster: my-first-cluster-1 Namespace: nginx-namespace Created: Apr 7, 2025, 11:08:00 PM Labels: app: nginx-deploy pod-template-hash: c9d9f6c6c Annotations: Not set Node: gke-my-first-cluster-1-default-pool-a492d1f-m4ww Termination grace period: 30 sec Restart policy: Always Phase: Running Start time: Apr 7, 2025, 11:08:00 PM Conditions: PodReadyToStartContainers: True Initialized: True Ready: True ContainerReady: True PodScheduled: True

Containers

Name	Status	Image	Restart count	Logs
nginx	Running	nginx	0	View logs

Exposing services

Name	Type	Endpoints
nginx-deploy-service	Load balancer	35.238.217.97:80



Google Cloud Suresh-Jenkins-Project Search (/) for resources, docs, products, and more Search Operations

Kubernetes Engine Learn about Enterprise Pod details Refresh Edit Delete Expose kubectl

nginx-deploy-c9d9f6c6c-dwt7m

Details Events Logs App Errors (0) YAML

Select the Cloud Monitoring account to see charts.

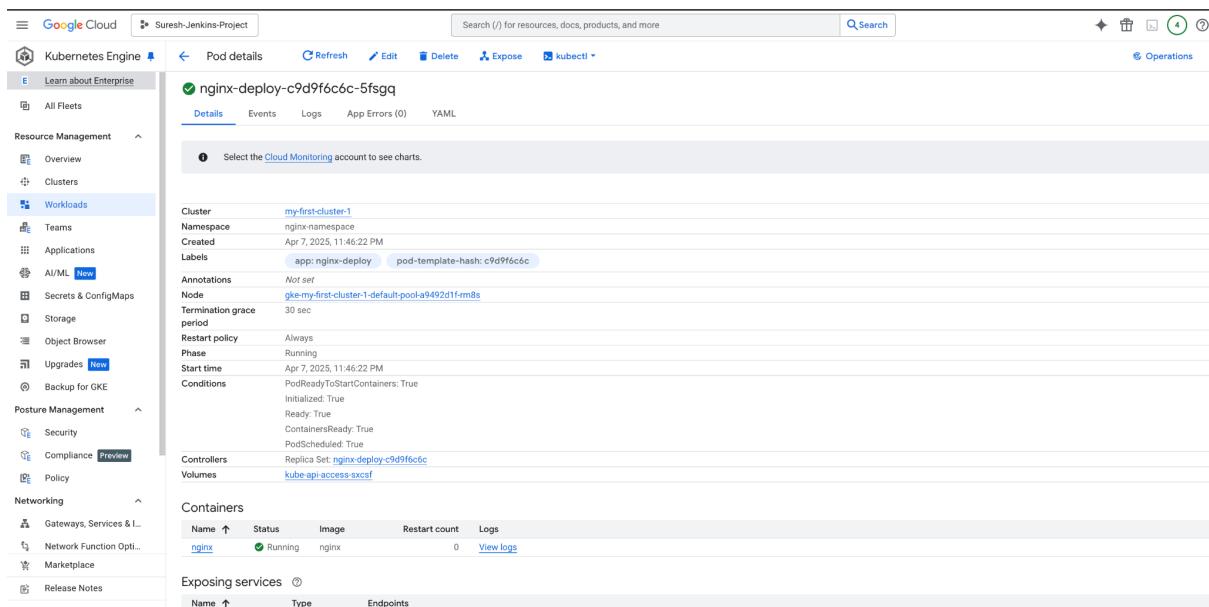
Cluster: my-first-cluster-1 Namespace: nginx-namespace Created: Apr 7, 2025, 11:15:46 PM Labels: app: nginx-deploy, pod-template-hash: c9d9f6c6c Annotations: Not set Node: gke-my-first-cluster-1-default-pool-a9492d1f-jhnx Termination grace period: 30 sec Restart policy: Always Phase: Running Start time: Apr 7, 2025, 11:15:46 PM Conditions: PodReadyToStartContainers: True, Initialized: True, Ready: True, ContainersReady: True, PodScheduled: True Controllers: Replica Set: nginx-deploy-c9d9f6c6c Volumes: kube-api-access-prwg

Containers

Name	Status	Image	Restart count	Logs
nginx	Running	nginx	0	View logs

Exposing services

Name	Type	Endpoints
nginx-deploy-service	Load balancer	35.238.217.97:80



Google Cloud Suresh-Jenkins-Project Search (/) for resources, docs, products, and more Search Operations

Kubernetes Engine Learn about Enterprise Pod details Refresh Edit Delete Expose kubectl

nginx-deploy-c9d9f6c6c-5fsgq

Details Events Logs App Errors (0) YAML

Select the Cloud Monitoring account to see charts.

Cluster: my-first-cluster-1 Namespace: nginx-namespace Created: Apr 7, 2025, 11:46:22 PM Labels: app: nginx-deploy, pod-template-hash: c9d9f6c6c Annotations: Not set Node: gke-my-first-cluster-1-default-pool-a9492d1f-rm8s Termination grace period: 30 sec Restart policy: Always Phase: Running Start time: Apr 7, 2025, 11:46:22 PM Conditions: PodReadyToStartContainers: True, Initialized: True, Ready: True, ContainersReady: True, PodScheduled: True Controllers: Replica Set: nginx-deploy-c9d9f6c6c Volumes: kube-api-access-sxcf

Containers

Name	Status	Image	Restart count	Logs
nginx	Running	nginx	0	View logs

Exposing services

Name	Type	Endpoints
nginx-deploy-service	Load balancer	35.238.217.97:80

✓ Task 5 – Set Up Horizontal Pod Autoscaler (HPA) in Kubernetes

Objective:

To deploy an Nginx application with CPU resource limits and configure a Horizontal Pod Autoscaler (HPA) to scale pods based on CPU usage.

Step-by-Step SOP

Step 1: Create the Nginx Deployment

Create a deployment named `nginx-deploy1` with the official Nginx image:

```
kubectl create deployment nginx-deploy1 --image=nginx
```

Step 2: Set CPU Resource Requests and Limits

Apply CPU resource limits to the deployment:

```
kubectl set resources deployment nginx-deploy1 --limits(cpu=200m) --requests(cpu=100m)
```

This ensures that the deployment is eligible for autoscaling based on CPU.

Step 3: Apply the Horizontal Pod Autoscaler (HPA)

Configure HPA to scale the deployment between 1 and 5 pods if CPU usage exceeds 50%:

```
kubectl autoscale deployment nginx-deploy1 --cpu-percent=50 --min=1 --max=5
```

Step 4: Verify the HPA Status

Check if the HPA has been created and is monitoring the deployment:

```
kubectl get hpa
```

```
ourbeastapi@cloudshell:~ (ashwin-451403)$ kubectl create deployment nginx-deploy2 --image=nginx
deployment.apps/nginx-deploy2 created
ourbeastapi@cloudshell:~ (ashwin-451403)$ kubectl set resources deployment nginx-deploy \
--limits(cpu=200m) \
--requests(cpu=100m)
Error from server (NotFound): deployments.apps "nginx-deploy" not found
ourbeastapi@cloudshell:~ (ashwin-451403)$ kubectl set resources deployment nginx-deploy2 --limits(cpu=200m) --requests(cpu=100m)
deployment.apps/nginx-deploy2 resource requirements updated
ourbeastapi@cloudshell:~ (ashwin-451403)$ kubectl autoscale deployment nginx-deploy2 \
--cpu-percent=50 \
--min=1 \
--max=5
horizontalpodautoscaler.autoscaling/nginx-deploy2 autoscaled
ourbeastapi@cloudshell:~ (ashwin-451403)$ kubectl get hpa
NAME          REFERENCE          TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
nginx-deploy2  Deployment/nginx-deploy2  cpu: <unknown>/50%  1          5          0          9s
ourbeastapi@cloudshell:~ (ashwin-451403)$ kubectl get svc
```

Thank You.

Suresh Dike.
